

Introduction to Bayesian linear regression with brms

Stefano Coretta

18/01/2020

Warning: Missing column names filled in: 'X46' [46]

Warning: Missing column names filled in: 'X50' [50]

Warning: Missing column names filled in: 'X50' [50]

Warning: Missing column names filled in: 'X50' [50]

- Safe method:
 - Install Rstan first: <https://github.com/stan-dev/rstan/wiki/RStan-Getting-Started> (see Installation of Rstan, Checking the C++ Toolchain, and Configuration of the C++ Toolchain).
 - Note that in Checking the C++ Toolchain details differ depending on OS.
 - Install brms: <https://github.com/paul-buerkner/brms#how-do-i-install-brms>.

Random variables

- We have a question about the world, so we collect data (sample from a population).
 - $y = (y_1, y_2, y_3, y_4, \dots, y_n)$
- We want to know how the data (the sample) was generated.
- In probability theory, data is generated by a random variable Y .

- Y is uncertain.
 - We can describe Y as a probability distribution, expressed by a set of parameters $\Theta = (\theta_1, \dots, \theta_n)$.
- Probability distributions:
 - *Normal* (μ, σ) ,
 - *Binomial* (n, p) ,
 - ...

$$vot_i \sim Normal(\mu, \sigma)$$

$$voiced_i \sim Bernoulli(p)$$

$$DoubleDative_i \sim Poisson(\lambda)$$

Frequentist vs Bayesian view

- Parameters: $\mu, \sigma, p, \lambda, \dots$
- Frequentist view:
 - The parameters are **fixed** (they are unknown but certain).
 - They take on a specific value.
- Bayesian view:
 - The parameters are **random variables** (they are unknown and uncertain).
 - We describe each parameter as a probability distribution, expressed by a set of **hyperparameters**.

$$vot_i \sim Normal(\mu, \sigma)$$

$$\mu \sim Normal(\mu_1, \sigma_1)$$

$$\sigma \sim HalfCauchy(x_0, \gamma)$$

Bayes' Theorem

$$P(\theta \mid d) = \frac{P(d \mid \theta) P(\theta)}{P(d)}$$


Bayes' Theorem

$$\textit{posterior probability} = \frac{\textit{likelihood} \times \textit{prior}}{\textit{marginal likelihood}}$$

- We can incorporate previous knowledge about the hyperparameters as **priors** (prior distributions).
- Priors are chosen based on expert knowledge, previous studies, pilot data...
 - Priors must **not** be chosen based on the data to be analysed.


- Informative and weakly informative priors.
- Uninformative or diffuse priors.
 - Uniform distribution.
- Regularising priors.


Normal prior



- Previous literature on VOT in Italian (Esposito, 2002; Stevens & Hajek, 2010) report VOT values for voiceless stops in the range of 20–60 ms.
 - We can express this knowledge with the prior $Normal(40, 10)$.
 - This is a somewhat strongly informative prior.

Normal(40, 10)






$$vot_i \sim Normal(\mu, \sigma)$$

$$\mu \sim Normal(40, 10)$$


$$\sigma \sim HalfCauchy(x_0, \gamma)$$

Cauchy prior



Cauchy prior

HalfCauchy



$$vot_i \sim Normal(\mu, \sigma)$$

$$\mu \sim Normal(40, 10)$$

$$\sigma \sim HalfCauchy(0, 10)$$

- We have a model which incorporates (some of) our knowledge about VOT (through the priors for μ and σ).
- Now we want to obtain the **posterior distributions** of μ and σ .
 - The posterior distribution is the prior distribution *conditioned* on the data.
- **brms** R package: Bayesian Regression Models using Stan (Bürkner, 2018).

- Stan (Stan Development Team, 2017).
 - Statistical programming language written in C++ for fitting Bayesian models (calculate posterior distributions).
 - Calculation can be complex and/or impossible, so we take many samples from the data and from the possible parameter values to find the posterior distributions of the hyperparameters.
 - Markov Chain Monte Carlo (MCMC) sampling using the No-U-Turn sampler (NUTS).
- brms is an interface between R and Stan.
- `brm()` function from brms.
 - lme4 syntax ($y \sim x + (1|w)$).
 - Creates a Stan model, which is compiled and run.

```
library(brms)

vot1 <- brm(
  <model_formula>,
  <family>,
  <prior>,
  <data>,
  chains = 4,
  iter = 2000
)
```

```
library(brms)

vot1 <- brm(
  vot ~ 1,
  family = gaussian(),
  <prior>,
  data = ita_egg,
  chains = 4,
  iter = 2000
)
```


Get prior

```
get_prior(  
  vot ~ 1,  
  family = gaussian(),  
  data = ita_egg  
)
```

```
##               prior      class coef group resp dpar nI  
## 1 student_t(3, 19, 14) Intercept  
## 2 student_t(3, 0, 14)      sigma
```


Prior predictive checks

```
nsim <- 1000
nobs <- 100

y <- matrix(rep(NA, nsim * nobs), ncol = nobs)


mu <- rnorm(nsim, 40, 10)
sigma <- rhcauchy(nsim, 10)

for (i in 1:nsim) {
  y[i,] <- rnorm(nobs, mean = mu[i], sd = sigma[i])
}
```


Prior predictive checks

```
op <- par(mfrow = c(3, 3))
rand_sample <- sample(1:nsim, 9)


for (i in rand_sample) {
  plot(density(y[i,]), main = "")
}
```




N = 100 Bandwidth = 7.379




N = 100 Bandwidth = 35.99




N = 100 Bandwidth = 1.575



N = 100 Bandwidth = 12.67



N = 100 Bandwidth = 4.789



N = 100 Bandwidth = 0.8624

Run the model

```
vot1 <- brm(  
  vot ~ 1,  
  family = gaussian(),  
  prior = priors,  
  data = ita_egg,  
  chains = 4,  
  iter = 2000,  
  file = "./cache/vot1"  
)
```

Model summary

vot1

Family: gaussian

Links: mu = identity; sigma = identity

Formula: vot ~ 1

Data: ita_egg (Number of observations: 2624)

Samples: 4 chains, each with iter = 2000; warmup = 1000

total post-warmup samples = 4000

##

Population-Level Effects:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk
--	----------	-----------	----------	----------	------	------

## Intercept	23.08	0.30	22.49	23.67	1.00	
--------------	-------	------	-------	-------	------	--

##

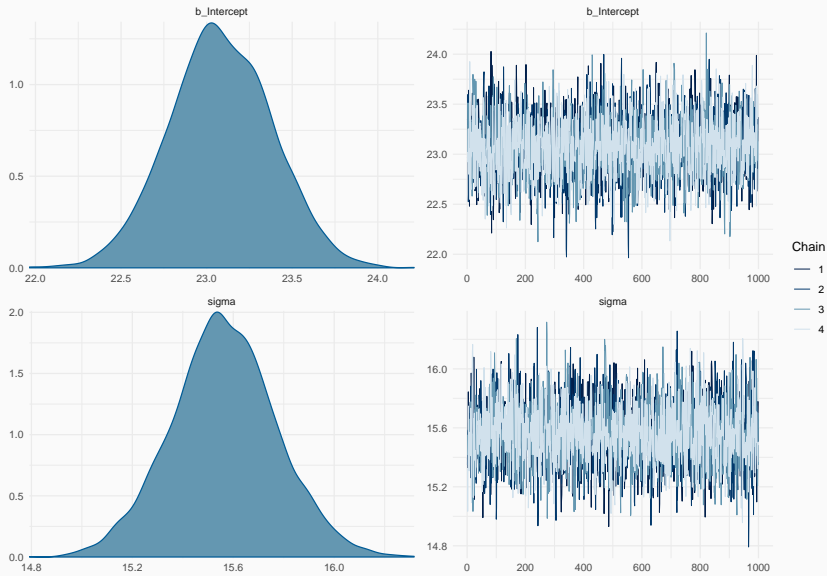
Family Specific Parameters:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk	ESS
--	----------	-----------	----------	----------	------	------	-----

## sigma	15.56	0.21	15.15	15.98	1.00		3391
----------	-------	------	-------	-------	------	--	------


Plot model

```
plot(vot1)
```



Posterior predictive check

```
pp_check(vot1, nsamples = 100)
```



Adding predictors

$$vot_i \sim Normal(\mu_i, \sigma)$$

$$\mu_i = \alpha + \beta_1 \times coronal_i + \beta_2 \times velar_i$$

$$\alpha \sim Normal(\mu_1, \sigma_1)$$

$$\beta_1 \sim Normal(\mu_2, \sigma_2)$$

$$\beta_2 \sim Normal(\mu_3, \sigma_3)$$

$$\sigma \sim HalfCauchy(x_0, \gamma)$$

Adding predictors

$$vot_i \sim Normal(\mu_i, \sigma)$$

$$\mu_i = \alpha + \beta_1 \times coronal_i + \beta_2 \times velar_i$$

$$\alpha \sim Normal(25, 10)$$

$$\beta_1 \sim Normal(10, 10)$$

$$\beta_2 \sim Normal(20, 10)$$

$$\sigma \sim HalfCauchy(0, 10)$$

Adding predictors

```
get_prior(  
  vot ~ 1 + c1_place,  
  family = gaussian(),  
  data = ita_egg  
)
```

```
##              prior      class      coef group 1  
## 1              b  
## 2              b c1_placecoronal  
## 3              b  c1_placevelar  
## 4 student_t(3, 19, 14) Intercept  
## 5 student_t(3, 0, 14)   sigma
```

Adding predictors

```
priors <- c(  
  prior(normal(25, 10), class = Intercept),  
  prior(cauchy(0, 10), class = sigma),  
  prior(normal(10, 10), class = b, coef = "c1_placecoronal"),  
  prior(normal(20, 10), class = b, coef = "c1_placevelar")  
)
```

Adding predictors

```
vot2 <- brm(  
  vot ~ 1 + c1_place,  
  family = gaussian(),  
  prior = priors,  
  data = ita_egg,  
  chains = 4,  
  iter = 2000,  
  file = "./cache/vot2"  
)
```

Random effects

$$vot_i \sim Normal(\mu_i, \sigma)$$

$$\mu_i =$$

$$\alpha + \alpha_{speaker[i]} + (\beta_1 + \beta_{1speaker[i]}) \times coronal_i + (\beta_2 + \beta_{2speaker[i]}) \times velar_i$$

$$\begin{bmatrix} \alpha_{speaker} \\ \beta_{1speaker[i]} \\ \beta_{2speaker[i]} \end{bmatrix} \sim MVNormal\left(\begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}, S\right)$$

$$\alpha \sim Normal(25, 10)$$

$$\alpha_{speaker} \sim Normal(0, \sigma_{speaker})$$

$$\beta_1 \sim Normal(10, 10)$$

$$\beta_2 \sim Normal(20, 10)$$

$$\sigma_{\alpha speaker} \sim Normal(0, \sigma_{speaker})$$

$$\sigma_{\beta 1 speaker} \sim HalfCauchy(0, 10)$$

Random effects

```
get_prior(  
  vot ~ 1 + c1_place + (1 + c1_place | speaker),  
  family = gaussian(),  
  data = ita_egg  
)
```

##		prior	class	coef	group
## 1			b		
## 2			b	c1_placecoronal	
## 3			b	c1_placevelar	
## 4		lkj(1)	cor		
## 5			cor		speaker
## 6	student_t(3, 19, 14)		Intercept		
## 7	student_t(3, 0, 14)		sd		
## 8			sd		speaker
## 9			sd	c1_placecoronal	speaker

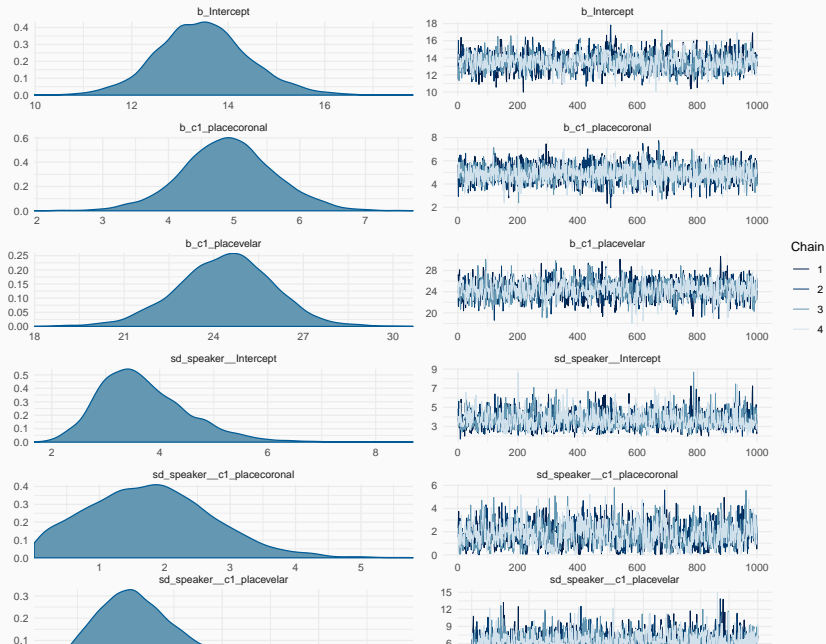
Random effects

```
priors <- c(  
  prior(normal(40, 10), class = Intercept),  
  prior(cauchy(0, 10), class = sigma),  
  prior(normal(10, 10), class = b, coef = "c1_placecoronal"),  
  prior(normal(20, 10), class = b, coef = "c1_placevelar"),  
  prior(normal(0, 25), class = sd),  
  prior(lkj(2), class = cor)  
)
```


Random effects

```
vot3 <- brm(  
  vot ~ 1 + c1_place + (1 + c1_place | speaker),  
  family = gaussian(),  
  prior = priors,  
  data = ita_egg,  
  chains = 4,  
  iter = 2000,  
  file = "./cache/vot3"  
)
```

Random effects



Random effects

vot3

```
## Family: gaussian
```

```
## Links: mu = identity; sigma = identity
```

```
## Formula: vot ~ 1 + c1_place + (1 + c1_place | speaker)
```

```
## Data: ita_egg (Number of observations: 2624)
```

```
## Samples: 4 chains, each with iter = 2000; warmup = 1000
```

```
## total post-warmup samples = 4000
```

```
##
```

```
## Group-Level Effects:
```

```
## ~speaker (Number of levels: 18)
```

```
##
```

	Estimate	Est.Error	1-
--	----------	-----------	----

```
## sd(Intercept)
```

3.70	0.83	
------	------	--

```
## sd(c1_placecoronal)
```

1.78	0.94	
------	------	--

```
## sd(c1_placevelar)
```

6.46	1.38	
------	------	--

```
## cor(Intercept,c1_placecoronal)
```

-0.24	0.32	
-------	------	--

Binomial logistic regression


```
priors <- c(  
  prior(student_t(3, 0, 2), class = Intercept),  
  prior(student_t(3, 0, 2), class = b),  
  prior(cauchy(0, 1), class = sd),  
  prior(lkj(2), class = cor)  
)
```

Binomial logistic regression

```
burst1 <- brm(  
  response_n ~  
    burst *  
    vowel +  
    (1+burst|participant),  
  data = burst,  
  prior = priors,  
  family = bernoulli,  
  file = "./cache/burst1",  
  control = list(adapt_delta = 0.999)  
)
```

Binomial logistic regression

```
conditional_effects(burst1, effects = "burst:vowel")
```






Figure 2: statistical inference

Open Learning Initiative, <https://oli.cmu.edu/courses/concepts-of-statistics/>
(CC BY-NC-SA 4.0)

- We want to know two things:
 - If there is **evidence for our hypothesis H** (or for the value of the parameter θ), and
 - What the **strength** of the evidence is.
- This is the aim of inferential statistics.
- We test H against empirical data (hypothesis testing).
 - It is important to decide in advance the details of the analysis (model and prior specification among other things).
- Inference is ultimately a long-term endeavour (via accumulation of knowledge).

- Three ways of doing inference (hypothesis testing) with Bayesian statistics:
 - Inference from the posterior.
 - Inference using a Region Of Practical Equivalence (ROPE).
 - Inference using the Bayes factor.

Inference from the posterior

- H: Condition B decreases reaction times.
- Posterior: 95% CI = $[-80, -15]$ ms.
- Inference: The posterior suggests that Condition B decreases reaction times by 15 to 80 ms at 95% confidence.

Inference with a ROPE

References

- Bürkner, Paul-Christian. 2018. Advanced Bayesian multilevel modeling with the R package brms. *The R Journal* 10(1). 395–411. doi:10.32614/RJ-2018-017.
- Esposito, Anna. 2002. On vowel height and consonantal voicing effects: Data from Italian. *Phonetica* 59(4). 197–231. doi:10.1159/000068347.
- Stan Development Team. 2017. Stan: A C++ library for probability and sampling, version 2.14.0. <http://mc-stan.org/>.

Stevens, Mary & John Hajek. 2010. Post-aspiration in standard Italian: some first cross-regional acoustic evidence. Paper presented at Interspeech, 26-30 September 2010, Makuhari, Chiba, Japan.